

# causal\_inference

December 5, 2022

## 1 Causal Inference

This notebook shows our work on our causal inference question: How does a mother's smoking behavior cause the change of her baby's birth weight?

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import scipy.stats as stats
from scipy.stats import beta, binom
import itertools
from ipywidgets import interact, interactive
import statsmodels.api as sm

import sklearn
from sklearn.linear_model import LogisticRegression as LR
```

```
[2]: #load dataset
birth = pd.read_csv('subsampled_clean_data.csv')
birth
```

```
[2]:
```

	ATTEND	BFACIL	BMI	DBWT	DMAR	FAGECOMB	FEDUC	FRACE6	LD_IND	MAGER	\
0	1	1	31.4	3670	1	29	6	1	N	32	
1	2	1	27.6	3494	1	34	4	1	Y	33	
2	1	1	27.1	3374	2	43	2	1	N	29	
3	1	1	26.8	3520	1	30	3	1	Y	28	
4	1	1	21.3	3140	1	30	5	1	N	30	
...	...	...	...	...	...	...	...	...	...	...	
9995	1	1	35.9	3062	1	30	4	1	Y	30	
9996	1	1	22.5	3855	1	30	3	1	Y	23	
9997	1	1	20.4	2710	1	39	2	1	Y	32	
9998	1	1	24.4	3118	1	35	2	1	Y	34	
9999	1	1	24.6	3020	1	32	3	1	N	28	

	...	PRIORLIVE	PRIORTERM	RDMETH_REC	RESTATUS	RF_CESAR	SEX	\
0	...	False	False	1	2	N	M	
1	...	True	False	1	1	N	F	

2	...	True	True	1	1	N	M
3	...	False	True	1	1	N	M
4	...	False	False	1	3	N	M
...	...	...	...	...	...	...	...
9995	...	True	False	1	1	N	M
9996	...	False	False	3	1	N	M
9997	...	True	True	1	2	N	M
9998	...	False	False	3	2	N	M
9999	...	True	False	3	1	N	M

	PREG_LEN	WTGAIN_PER	CIG	FIRST_BIRTH
0	9	0.000000	False	True
1	9	0.120482	False	False
2	10	0.061350	True	False
3	9	0.301282	False	True
4	9	0.208333	False	True
...	...	...	...	...
9995	9	0.173684	False	False
9996	9	0.263514	False	True
9997	9	0.388889	True	False
9998	10	0.147887	False	True
9999	9	0.192308	True	False

[10000 rows x 31 columns]

## 1.1 Causal Inference - Randomized Experiments

Although the data is from an observational study, we could try to assume it is a randomized experiment and conduct causal inference. It can be served as a comparison with the later causal inference results when we use observational study techniques.

For randomized experiment technique, we use the Fisher Randomized Test (*i.e.*, Permutation Test) with the simple difference in means ( $\hat{\tau} = \frac{1}{n_1} \sum_{i=1}^n Z_i Y_i - \frac{1}{n_0} \sum_{i=1}^n (1 - Z_i) Y_i$ ) as the test statistic. The null hypothesis we are testing is

$$H_0 : Y_i(1) = Y_i(0) \quad \forall i = 1, \dots, n$$

which is also known as *sharp/strong null hypothesis*. It is basically saying that the treatment and control outcomes come from the *same* distribution.

First we compute the observed test statistic.

```
[3]: observed_T = np.mean(birth[birth["CIG"] == True].DBWT) - np.
      ↪ mean(birth[birth["CIG"] == False].DBWT)
      observed_T
```

```
[3]: -115.41768336628411
```

```
[4]: birth['CIG'].sum()
```

```
[4]: 682
```

Since we have 10000 units with 682 treated units, going through all possible permutations – there are  $\binom{10000}{682}$  different permutations) – will be too time-consuming. We can use Monte Carlo to approximate the true p-value.

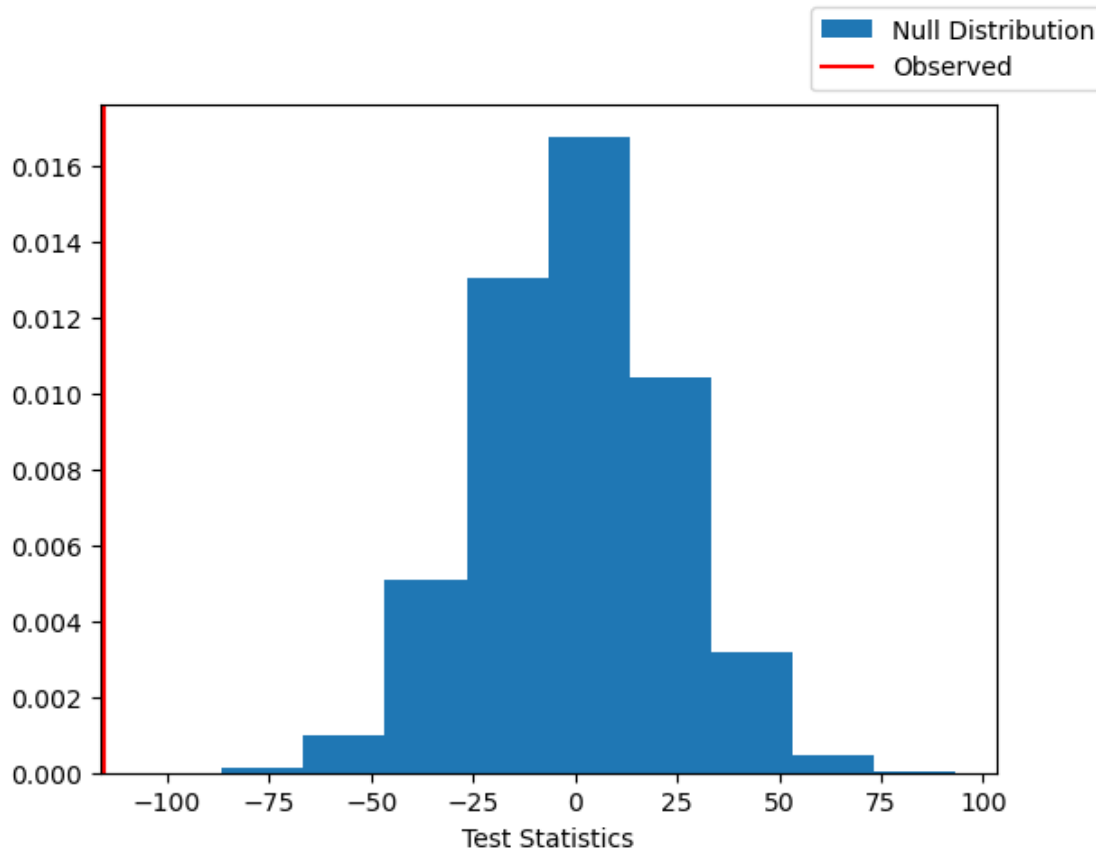
```
[5]: rng = np.random.default_rng(102)
R = 50000 # repetition times
Ts = np.zeros(R)
shuffled_birth = birth.copy()

for i in range(R):
    shuffled_birth['shuffled_CIG'] = rng.choice(birth['CIG'], size=birth.
↪shape[0], replace=False)
    Ts[i] = np.mean(shuffled_birth[shuffled_birth["shuffled_CIG"] == True].
↪DBWT) - np.mean(shuffled_birth[shuffled_birth["shuffled_CIG"] == False].DBWT)

p_val = np.sum(np.abs(Ts) >= np.abs(observed_T)) / R
print(f'The p-value is {p_val}')

fig, ax = plt.subplots()
ax.hist(Ts, density=True, label='Null Distribution')
ax.axvline(observed_T, color='r', label='Observed')
ax.set_xlabel('Test Statistics')
fig.legend()
fig.show()
```

The p-value is 0.0



Based on the approximated p-value, it is clear that we should reject the null hypothesis. Thus, we claim that a mother's smoking behavior will cause the change of her baby's birth weight. However, the permutation test cannot tell either qualitative (*i.e.*, whether smoking causes lower or higher birth weight) or quantitative (*i.e.*, how much birth weight increase or decrease can be caused by smoking) causal effect. To do so, we use the observational study techniques, and we will explain them next.

## 1.2 Causal Inference - Observational Studies

```
[6]: causal_effect = np.mean(birth[birth["CIG"] == True].DBWT) - np.
    ↪ mean(birth[birth["CIG"] == False].DBWT)
    causal_effect
```

```
[6]: -115.41768336628411
```

Using the simple difference in means, the causal effect of cigarettes on baby's birth weight is negative, meaning that smoking will lead to a lower birth weight.

### 1.3 Outcome Regression

We know that BMI and prenatal care may be confounders. Then by unconfoundedness, we could fit a linear model of the following form:

$$\text{Birth Weights} = \alpha Z + \beta \text{BMI} + \gamma \text{PRECARE}$$

If we make two assumptions, then the estimated coefficient of treatment from OLS,  $\hat{\tau}$ , will be an unbiased estimate of the ATE. The two assumptions are:

1. Assume unconfoundedness given BMI and PRECARE.
2. Assume this linear model correctly describes the interaction between the variables.

First, in order to fit the regression, we need to change categorical variable CIG to dummy variable CIG\_True.

```
[7]: #change categorical data to dummy variable
birth = pd.get_dummies(birth, columns=['CIG'], drop_first=True)
birth
```

```
[7]:
```

	ATTEND	BFACIL	BMI	DBWT	DMAR	FAGECOMB	FEDUC	FRACE6	LD_IND	MAGER	\
0	1	1	31.4	3670	1	29	6	1	N	32	
1	2	1	27.6	3494	1	34	4	1	Y	33	
2	1	1	27.1	3374	2	43	2	1	N	29	
3	1	1	26.8	3520	1	30	3	1	Y	28	
4	1	1	21.3	3140	1	30	5	1	N	30	
...	...	...	...	...	...	...	...	...	...	...	
9995	1	1	35.9	3062	1	30	4	1	Y	30	
9996	1	1	22.5	3855	1	30	3	1	Y	23	
9997	1	1	20.4	2710	1	39	2	1	Y	32	
9998	1	1	24.4	3118	1	35	2	1	Y	34	
9999	1	1	24.6	3020	1	32	3	1	N	28	

	...	PRIORLIVE	PRIORTERM	RDMETH_REC	RESTATUS	RF_CESAR	SEX	\
0	...	False	False	1	2	N	M	
1	...	True	False	1	1	N	F	
2	...	True	True	1	1	N	M	
3	...	False	True	1	1	N	M	
4	...	False	False	1	3	N	M	
...	...	...	...	...	...	...	...	
9995	...	True	False	1	1	N	M	
9996	...	False	False	3	1	N	M	
9997	...	True	True	1	2	N	M	
9998	...	False	False	3	2	N	M	
9999	...	True	False	3	1	N	M	

	PREG_LEN	WTGAIN_PER	FIRST_BIRTH	CIG_True
0	9	0.000000	True	0
1	9	0.120482	False	0
2	10	0.061350	False	1

3	9	0.301282	True	0
4	9	0.208333	True	0
...	...	...	...	...
9995	9	0.173684	False	0
9996	9	0.263514	True	0
9997	9	0.388889	False	1
9998	10	0.147887	True	0
9999	9	0.192308	False	1

[10000 rows x 31 columns]

treatment (Z): birth['CIG\_True']

outcome (Y): birth['DBWT']

confounder (X): birth['BMI'], birth['PRECARE']

units: baby's birth weights

```
[8]: def fit_OLS_model(df, target_variable, explanatory_variables, intercept =
      False):
      target = df[target_variable]
      inputs = df[explanatory_variables]
      if intercept:
          inputs = sm.add_constant(inputs)

      fitted_model = sm.OLS(target, inputs).fit()
      return(fitted_model)

      def mean_squared_error(true_vals, predicted_vals):
          return np.mean((true_vals - predicted_vals) ** 2)

[9]: full_linear_model = fit_OLS_model(birth, 'DBWT', ['CIG_True', 'BMI', 'PRECARE'])
      print(full_linear_model.summary())
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          DBWT    R-squared (uncentered):
0.933
Model:                  OLS     Adj. R-squared (uncentered):
0.933
Method:                 Least Squares    F-statistic:
4.611e+04
Date:                   Mon, 05 Dec 2022    Prob (F-statistic):
0.00
Time:                   21:50:04    Log-Likelihood:
-81848.
No. Observations:      10000    AIC:
```

```

1.637e+05
Df Residuals:          9997    BIC:
1.637e+05
Df Model:              3
Covariance Type:      nonrobust
=====
              coef    std err          t      P>|t|      [0.025    0.975]
-----
CIG_True    -104.8087     34.480     -3.040     0.002    -172.397    -37.221
BMI           88.6601      0.707    125.335     0.000      87.273     90.047
PRECARE      642.9650     15.157     42.421     0.000     613.254     672.675
=====
Omnibus:                 1056.169    Durbin-Watson:                 1.978
Prob(Omnibus):              0.000    Jarque-Bera (JB):             1570.134
Skew:                      -0.797    Prob(JB):                     0.00
Kurtosis:                   4.107    Cond. No.                     111.
=====

```


Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[10]: sns.histplot(data=birth, x='DBWT', hue='CIG_True', stat='density')
      plt.title("Distributions of Babies' Weights under Different Cigarettes Usage")
```

```
[10]: Text(0.5, 1.0, "Distributions of Babies' Weights under Different Cigarettes Usage")
```




```
[11]: birth.groupby('CIG_True')['DBWT'].mean()
```

```
[11]: CIG_True
0    3299.782786
1    3184.365103
Name: DBWT, dtype: float64
```

```
[12]: sns.histplot(data=birth, x='DBWT', hue='PRECARE', stat='density')
plt.title("Distributions of Babies' Weights under Different Prenatal Care_
↪Levels")
```

```
[12]: Text(0.5, 1.0, "Distributions of Babies' Weights under Different Prenatal Care
Levels")
```





```
[13]: birth.groupby('PRECARE')['DBWT'].mean()
```

```
[13]: PRECARE
0    3056.591398
1    3301.014548
2    3253.022464
3    3295.043228
Name: DBWT, dtype: float64
```

From the outcome regression results table, we can see that by using BMI and PRECARE as confounders, the causal coefficient of cigarettes consumption on babies' weights is around -104.8. This means that smoking an extra cigarette will likely decrease the baby's weight by 104.8 grams.

From the graphs above, we can also see that the distribution of babies' weights, whether mother smoke or not, are uniformly distributed. However, babies whose mothers don't smoke tend to weigh more than those whose mothers do, having a difference in mean around 115 grams.

We can see that the distribution of babies' weights with different levels of prenatal care are also uniformly distributed. Generally, more prenatal care will lead to higher babies weights, and level-3 precare has the highest average weight of 3295 grams. However, we can see that the average weights of babies under level-1 precare is actually higher for those under level-2. What causes this difference is unknown.

### 1.3.1 Inverse Propensity Score

In this section, we use inverse propensity weighting.

Propensity score calculates the probability that a unit was treated, conditioned on a particular set of confounders  $x$ :

$$e(x) = P(Z = 1|X = x)$$

For inverse propensity weighting (IPW). Let  $n = n_0 + n_1$  be the total number of observations. The IPW estimator of the ATE is:

$$\hat{\tau}_{IPW} = \frac{1}{n} \sum_{i:Z_i=1} \frac{Y_i}{e(X_i)} - \frac{1}{n} \sum_{i:Z_i=0} \frac{Y_i}{1 - e(X_i)}$$

```
[14]: Z = birth.CIG_True.values
      Y = birth.DBWT.values
      X = birth[['BMI', 'PRECARE']]

      lr = LR(max_iter=200, random_state=0)
      lr.fit(X, Z)
```

```
[14]: LogisticRegression(max_iter=200, random_state=0)
```

```
[15]: birth['pscore'] = lr.predict_proba(X)[: ,1]
      birth
```

```
[15]:
```

	ATTEND	BFACIL	BMI	DBWT	DMAR	FAGECOMB	FEDUC	FRACE6	LD_IND	MAGER	\
0	1	1	31.4	3670	1	29	6	1	N	32	
1	2	1	27.6	3494	1	34	4	1	Y	33	
2	1	1	27.1	3374	2	43	2	1	N	29	
3	1	1	26.8	3520	1	30	3	1	Y	28	
4	1	1	21.3	3140	1	30	5	1	N	30	
...	...	...	...	...	...	...	...	...	...	...	
9995	1	1	35.9	3062	1	30	4	1	Y	30	
9996	1	1	22.5	3855	1	30	3	1	Y	23	
9997	1	1	20.4	2710	1	39	2	1	Y	32	
9998	1	1	24.4	3118	1	35	2	1	Y	34	
9999	1	1	24.6	3020	1	32	3	1	N	28	

	...	PRIORTERM	RDMETH_REC	RESTATUS	RF_CESAR	SEX	PREG_LEN	\
0	...	False	1	2	N	M	9	
1	...	False	1	1	N	F	9	
2	...	True	1	1	N	M	10	
3	...	True	1	1	N	M	9	
4	...	False	1	3	N	M	9	
...	...	...	...	...	...	...	...	

9995	...	False	1	1	N	M	9
9996	...	False	3	1	N	M	9
9997	...	True	1	2	N	M	9
9998	...	False	3	2	N	M	10
9999	...	False	3	1	N	M	9

	WTGAIN_PER	FIRST_BIRTH	CIG_True	pscore
0	0.000000	True	0	0.068063
1	0.120482	False	0	0.063830
2	0.061350	False	1	0.063292
3	0.301282	True	0	0.062971
4	0.208333	True	0	0.057350
...	...	...	...	...
9995	0.173684	False	0	0.073410
9996	0.263514	True	0	0.058535
9997	0.388889	False	1	0.076169
9998	0.147887	True	0	0.060458
9999	0.192308	False	1	0.081689

[10000 rows x 32 columns]

```
[16]: n = len(birth)
ipw = np.sum(birth[birth["CIG_True"] == 1].DBWT / birth[birth["CIG_True"] == 1].
        ↳pscore)/n - np.sum(birth[birth["CIG_True"] == 0].DBWT / (1 -
        ↳birth[birth["CIG_True"] == 0].pscore))/n
ipw
```

[16]: -121.73829642696728

We can see that by applying inverse propensity weighting, the coefficient of cigarettes on babies' birth is around -121.74, meaning that smoking an extra cigarette will likely decrease the baby's weight by 121.74 grams.

Anomalies happens if some observations are rare in the treatment group (i.e.  $( ) 0$ ), which may cause the inverse propensity score,  $1/( )$  to be enormous. Therefore, we decide to only include points with propensity scores between 0.1 and 0.9 which accepts some bias to reduce the variance.

```
[17]: birth_new = birth[(birth['pscore'] > 0.1) & (birth['pscore'] < 0.9)]
n = len(birth_new)
trimmed_ipw = np.sum(birth_new[birth_new["CIG_True"] == 1].DBWT /
        ↳birth_new[birth_new["CIG_True"] == 1].pscore)/n - np.
        ↳sum(birth_new[birth_new["CIG_True"] == 0].DBWT / (1 -
        ↳birth_new[birth_new["CIG_True"] == 0].pscore))/n
trimmed_ipw
```

[17]: -593.0929729236877

Now, we can see that the coefficient of cigarettes on babies' birth becomes larger, around -593,

meaning that smoking an extra cigarette will likely decrease the baby's weight by 593 grams.

# prediction\_glm

December 5, 2022

## 1 Prediction

This notebook shows our work on our prediction question: How to predict the birth weight?

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
```

```
[2]: birth = pd.read_csv('subsampled_clean_data.csv')
```

```
[3]: birth.head()
```

```
[3]:   ATTEND  BFACIL   BMI  DBWT  DMAR  FAGECOMB  FEDUC  FRACE6  LD_INDL  MAGER  \
0        1        1  31.4  3670     1         29     6         1         N     32
1        2        1  27.6  3494     1         34     4         1         Y     33
2        1        1  27.1  3374     2         43     2         1         N     29
3        1        1  26.8  3520     1         30     3         1         Y     28
4        1        1  21.3  3140     1         30     5         1         N     30
```

```
   ...  PRIORLIVE  PRIORTERM  RDMETH_REC  RESTATUS  RF_CESAR  SEX  PREG_LEN  \
0  ...        False        False           1         2         N    M         9
1  ...         True        False           1         1         N    F         9
2  ...         True         True           1         1         N    M        10
3  ...        False         True           1         1         N    M         9
4  ...        False        False           1         3         N    M         9
```

```
   WTGAIN_PER  CIG  FIRST_BIRTH
0    0.000000  False         True
1    0.120482  False        False
2    0.061350   True        False
3    0.301282  False         True
4    0.208333  False         True
```

```
[5 rows x 31 columns]
```

## 2 Research Question

Predicting a baby's birth weight from Mother's Single Years of Age and Number of Prenatal Visits, comparing GLMs to nonparametric methods.

## 3 Goal

We are trying to use various models including GLMs and Nonparametric methods to predict baby's birth weights. In the original EDA and using our domain knowledge on newborn health, we proposed two features that will be helpful for constructing the model, **Mother's Single Years of Age (MAGER)** and **Number of Prenatal Visits (PREVIS)**.

**MAGER:** Using our common knowledge, we think that younger mothers are likely to have healthier babies and thus higher birth weights.

**PREVIS:** Using our common knowledge, we think that more prenatal visits means that the family pays more attention to the pregnancy and thus is likely to have babies with higher birth weights.

## 4 Nonparametric Method-Decision Tree

### Why Decision Tree

For nonparametric method like decision tree, we do not need to consider about selecting the features. Since the impurity and impurity reduction will do the heavy lifting. The first few depths will select the best features and threshold that decrease the impurity the most, which help us to see which features are the best for the prediction.

**Assumption 1.** By using Decision Tree, we assume that the relationship is non linear and complex. 2. Since y is continuous, we will be using regression tree from CART, the Decision-TreeRegressor. We need to realize some mechanics of the model in order to make some assumption about our dataset: The model takes in X and y as input, and it tries to iterate through all the possible features in X and iterate through all the values(threshold) that particular feature can be. And calculate the impurity reduction of such split, where

$$Impurity = \frac{1}{N_n(t)} \sum_{i: X_i \in R_i} (y_i - \mu_n(t))^2 \quad (1)$$

and **Impurity reduction**

$$\Delta_I(t) = Impurity(t) - \frac{N_n(t^{left})}{N_n(t)} Impurity(t^{left}) - \frac{N_n(t^{right})}{N_n(t)} Impurity(t^{right}) \quad (2)$$

Impurity calculates the weighted sum of difference between mean and each sample's label squared in a tree node, which basically tells how pure a particular node is. The Impurity reduction calculates the difference between the Impurity before split and the weighted sum of the Impurity of nodes after split. Since there's mean involved in the model, we need to make all the data to continuous, we will achieve this by performing one-hot encoding.

#### 4.0.1 Baby Weights (y)

```
[4]: y = np.array(birth['DBWT'])
      y
```

```
[4]: array([3670, 3494, 3374, ..., 2710, 3118, 3020])
```

#### 4.0.2 Valid features to construct the tree (X)

```
[5]: # Drop irrelevant features and y
      X_drop = birth.drop(['DBWT', 'DMAR'], axis = 1)
      X_drop.head()
```

```
[5]:
```

	ATTEND	BFACIL	BMI	FAGECOMB	FEDUC	FRACE6	LD_INDL	MAGER	MBSTATE_REC	\
0	1	1	31.4	29	6	1	N	32	1	
1	2	1	27.6	34	4	1	Y	33	1	
2	1	1	27.1	43	2	1	N	29	1	
3	1	1	26.8	30	3	1	Y	28	1	
4	1	1	21.3	30	5	1	N	30	1	

	MEDUC	...	PRIORLIVE	PRIORTERM	RDMETH_REC	RESTATUS	RF_CESAR	SEX	\
0	6	...	False	False	1	2	N	M	
1	7	...	True	False	1	1	N	F	
2	2	...	True	True	1	1	N	M	
3	7	...	False	True	1	1	N	M	
4	4	...	False	False	1	3	N	M	

	PREG_LEN	WTGAIN_PER	CIG	FIRST_BIRTH
0	9	0.000000	False	True
1	9	0.120482	False	False
2	10	0.061350	True	False
3	9	0.301282	False	True
4	9	0.208333	False	True

[5 rows x 29 columns]

```
[6]: # Get all the categorical data
      cat_cols = X_drop.select_dtypes(exclude=["number"]).columns
      cat_cols
```

```
[6]: Index(['LD_INDL', 'PRIORDEAD', 'PRIORLIVE', 'PRIORTERM', 'RF_CESAR', 'SEX',
          'CIG', 'FIRST_BIRTH'],
          dtype='object')
```

```
[7]: for c in cat_cols:
      encoded = pd.get_dummies(X_drop[c], prefix=c)
      X_drop = pd.concat([X_drop, encoded], axis='columns')
```

```
[8]: X_encoded = X_drop.drop(cat_cols, axis = 1)
```

```
[9]: X_encoded.head()
```

```
[9]:
```

	ATTEND	BFACIL	BMI	FAGECOMB	FEDUC	FRACE6	MAGER	MBSTATE_REC	MEDUC	\
0	1	1	31.4	29	6	1	32	1	6	
1	2	1	27.6	34	4	1	33	1	7	
2	1	1	27.1	43	2	1	29	1	2	
3	1	1	26.8	30	3	1	28	1	7	
4	1	1	21.3	30	5	1	30	1	4	

	MRAVE6	...	PRIORTERM_False	PRIORTERM_True	RF_CESAR_N	RF_CESAR_Y	\
0	1	...	1	0	1	0	
1	1	...	1	0	1	0	
2	1	...	0	1	1	0	
3	1	...	0	1	1	0	
4	1	...	1	0	1	0	

	SEX_F	SEX_M	CIG_False	CIG_True	FIRST_BIRTH_False	FIRST_BIRTH_True	\
0	0	1	1	0	0	1	
1	1	0	1	0	1	0	
2	0	1	0	1	1	0	
3	0	1	1	0	0	1	
4	0	1	1	0	0	1	

[5 rows x 37 columns]

#### 4.0.3 Fit

```
[10]: X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.
↳01, random_state=0)
```

```
[11]: depths = [i for i in range(1, 10)]
train_scores = np.ones(len(depths))
test_scores = np.ones(len(depths))
```

#### Cross validation

```
[12]: for idx in range(len(depths)):
    clf = tree.DecisionTreeRegressor(max_depth = depths[idx])
    clf.fit(X_train, y_train)
    train_scores[idx] = clf.score(X_train, y_train)
    test_scores[idx] = clf.score(X_test, y_test)
    print("Max depths" ,depths[idx] , "will have train score" ,
↳train_scores[idx] , "and test score" , test_scores[idx])
```

Max depths 1 will have train score 0.15722212437225924 and test score 0.12774532540632055



Max depths 2 will have train score 0.23653748245476214 and test score 0.22971336641048823  
Max depths 3 will have train score 0.26145146321359913 and test score 0.21841821841646225  
Max depths 4 will have train score 0.2864851728757737 and test score 0.2692236761573288  
Max depths 5 will have train score 0.30956300362718303 and test score 0.2591029708037196  
Max depths 6 will have train score 0.33911101131747656 and test score 0.26877151389971643  
Max depths 7 will have train score 0.3766485557353514 and test score 0.2162482576076865  
Max depths 8 will have train score 0.42306787845459326 and test score 0.2156016222426992  
Max depths 9 will have train score 0.47492748583367006 and test score 0.15871008237693107

According to the cross validation, `max_depth = 4` looks promising, so that's what we will use in the prediction phase

```
[13]: clf = tree.DecisionTreeRegressor(max_depth = 4)
      clf.fit(X_train, y_train)
```

```
[13]: DecisionTreeRegressor(max_depth=4)
```


#### 4.0.4 Summarize and interpret

##### Visualize tree

```
[14]: plt.figure(figsize=(10,10))
      tree.plot_tree(clf, fontsize=9)
      plt.show()
```



```
plt.ylabel("Baby's birth weights")
plt.xlabel("Testing sample data")
plt.legend()
plt.show()
```



As the plot indicates, although the score of the model isn't the best, but it actually perform decently. The overlap between true label and predicted label is fair considering they are continuous.

## 5 Nonparametric Method-Random Forest

```
[17]: clf1 = RandomForestRegressor(n_estimators= 200,max_depth=15)
```

```
[18]: clf1.fit(X_train, y_train)
```

```
[18]: RandomForestRegressor(max_depth=15, n_estimators=200)
```

```
[19]: clf1.score(X_train, y_train)
```

```
[19]: 0.7983857863485958
```

```
[20]: clf1.score(X_test, y_test)
```


```
[20]: 0.32263016813649614
```

### 5.0.1 Summarize and interpret

We are unable to interpret the model since random forest is an ensemble algorithm that takes vote from hundreds of learners. Although the crowd's decision helps to lower the variance, it is unable to interpret anymore

#### Plotting the true labels and predicted labels

```
[21]: plt.rcParams['figure.figsize'] = (10, 6)
x_ax = range(len(X_test))
plt.plot(x_ax, y_test, label = 'true label', color = 'k', linestyle = '-')
plt.plot(x_ax, clf1.predict(X_test), label = 'predicted', color = 'k',
        linestyle = '--')
plt.ylabel("Baby's birth weights")
plt.xlabel("Testing sample data")
plt.legend()
plt.show()
```



## 6 GLM

```
[22]: #import the pymc3 package
import pymc3 as pm
from pymc3 import glm
import statsmodels.api as sm
import arviz as az
```

## 6.1 Choice of Model - Linear Regression

From EDA, we see that DBWT, MAGER, and PREVIS plots are all roughly normal distributed, and we are predicting real-valued outputs, so the best choice of model here is Linear Regression.

**Inverse Link Function:** Identity

**Likelihood:** Gaussian

This means that we will model birth weight as  $W_i \sim N(\beta_0 + \beta_1 M_i + \beta_2 P_i, \sigma^2 I_n)$  where  $M_i$  is MAGER and  $P_i$  is PREVIS.

## 6.2 Assumptions

For our model (Linear Regression), we are making the following assumptions: 1. There is a linear relationship between the response variable(DBWT) and explanatory variables (MAGER and PREVIS) 2. Constant Variance 3. Birth weights are assumed to be normally distributed (The histogram from EDA shows that this is valid). 4. The Birth weights are independently distributed.

## 6.3 Frequentist Regression

```
[23]: freq_model = sm.GLM(birth.DBWT, exog = sm.  
    ↪add_constant(birth[['MAGER', 'PREVIS']]),  
        family=sm.families.Gaussian())  
freq_res = freq_model.fit()  
print(freq_res.summary())
```

### Generalized Linear Model Regression Results

=====						
Dep. Variable:		DBWT	No. Observations:		10000	
Model:		GLM	Df Residuals:		9997	
Model Family:		Gaussian	Df Model:		2	
Link Function:		identity	Scale:		3.2879e+05	
Method:		IRLS	Log-Likelihood:		-77704.	
Date:	Mon, 05 Dec 2022	Deviance:		3.2869e+09		
Time:	21:48:09	Pearson chi2:		3.29e+09		
No. Iterations:		3				
Covariance Type:		nonrobust				
=====						
	coef	std err	z	P> z	[0.025	0.975]
-----						
const	3014.8158	34.136	88.318	0.000	2947.911	3081.721
MAGER	2.0604	1.031	1.998	0.046	0.039	4.081
PREVIS	18.5685	1.454	12.775	0.000	15.720	21.417
=====						

```
/opt/conda/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142:
```

```
FutureWarning: In a future version of pandas all arguments of concat except for  
the argument 'objs' will be keyword-only
```

```
    x = pd.concat(x[:, :order], 1)
```


We see that the prediction is  $W_i = 3014.8158 + 2.0604M_i + 18.5685P_i$ . This means that older mothers have babies with higher birth weights, and mothers who go to more prenatal visits also have babies with higher birth weights.

### 6.3.1 Model Checking

```
[24]: from statsmodels.graphics.api import abline_plot
```

```
[25]: nobs = freq_res.nobs  
y = np.array(birth['DBWT'])  
yhat = freq_res.mu
```

```
[26]: fig, ax = plt.subplots()  
ax.scatter(yhat, y)  
line_fit = sm.OLS(y, sm.add_constant(yhat, prepend=True)).fit()  
abline_plot(model_results=line_fit, ax=ax)  
  
ax.set_title('Model Fit Plot')  
ax.set_ylabel('Observed values')  
ax.set_xlabel('Fitted values');
```




From the model fit plot above, the prediction it provides is fine but there's still a lot of room for improvement, so we might want to try a different set of features to see if it gives better predictions (For example, using 'PREG\_LEN' and 'M\_Ht\_In', which are the best features from the decision


tree).

```
[27]: from statsmodels import graphics
      resid = freq_res.resid_deviance.copy()
      graphics.gofplots.qqplot(resid, line='r')
```

```
/opt/conda/lib/python3.9/site-packages/statsmodels/graphics/gofplots.py:993:
UserWarning: marker is redundantly defined by the 'marker' keyword argument and
the fmt string "bo" (-> marker='o'). The keyword argument will take precedence.
      ax.plot(x, y, fmt, **plot_style)
```

[27]:





The Q-Q Plots is roughly linear, which means that the assumption that the residuals are normally distributed is satisfied.

### 6.3.2 Uncertainty Quantification

The parameter we are estimating is a fixed quantity but our estimate is random, because it depends on our data and the data is random. According to the frequentist model, the 95% confidence interval for the intercept is [2947.911, 3081.721], for the coefficient of MAGER is [0.039, 4.081], for the coefficient of PREVIS is [15.720, 21.417]. The standard error for the intercept is 34.136, for the coefficient of MAGER is 1.031, for the coefficient of PREVIS is 1.454.

## 6.4 Bayesian Regression

```
[28]: dbwt = np.array(birth['DBWT'])
      mager = np.array(birth['MAGER'])
      previs = np.array(birth['PREVIS'])
```

### Choice of Priors:

For standard deviance, we set the prior to be exponential(0.01), because this must be nonnegative so exponential distribution is a good fit; according to the histogram we plotted in EDA, the SD of the birth weights is about 6, so the parameter 0.01 makes sense. For intercept and coefficients, we reference the result from the frequentist model and set relatively larger variances so that the posterior will be more dependent on the data instead of the priors.

```
[29]: with pm.Model() as bayes_model:
      #define the priors
```



```

sigma = pm.Exponential('sigma', lam=0.01)
intercept = pm.Normal("Intercept", 3015, sigma=30)
beta_1 = pm.Normal("beta_1", 2, sigma=3)
beta_2 = pm.Normal("beta_2", 18, sigma=3)

#likelihood
likelihood = pm.Normal("y", mu = intercept + beta_1*mager + beta_2*previs,
↳sigma = sigma, observed = dbwt)

#inference
trace = pm.sample(1000, cores = 2, target_accept = 0.95,
↳return_inferencedata=True)

```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt\_diag...

Multiprocess sampling (2 chains in 2 jobs)

NUTS: [beta\_2, beta\_1, Intercept, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

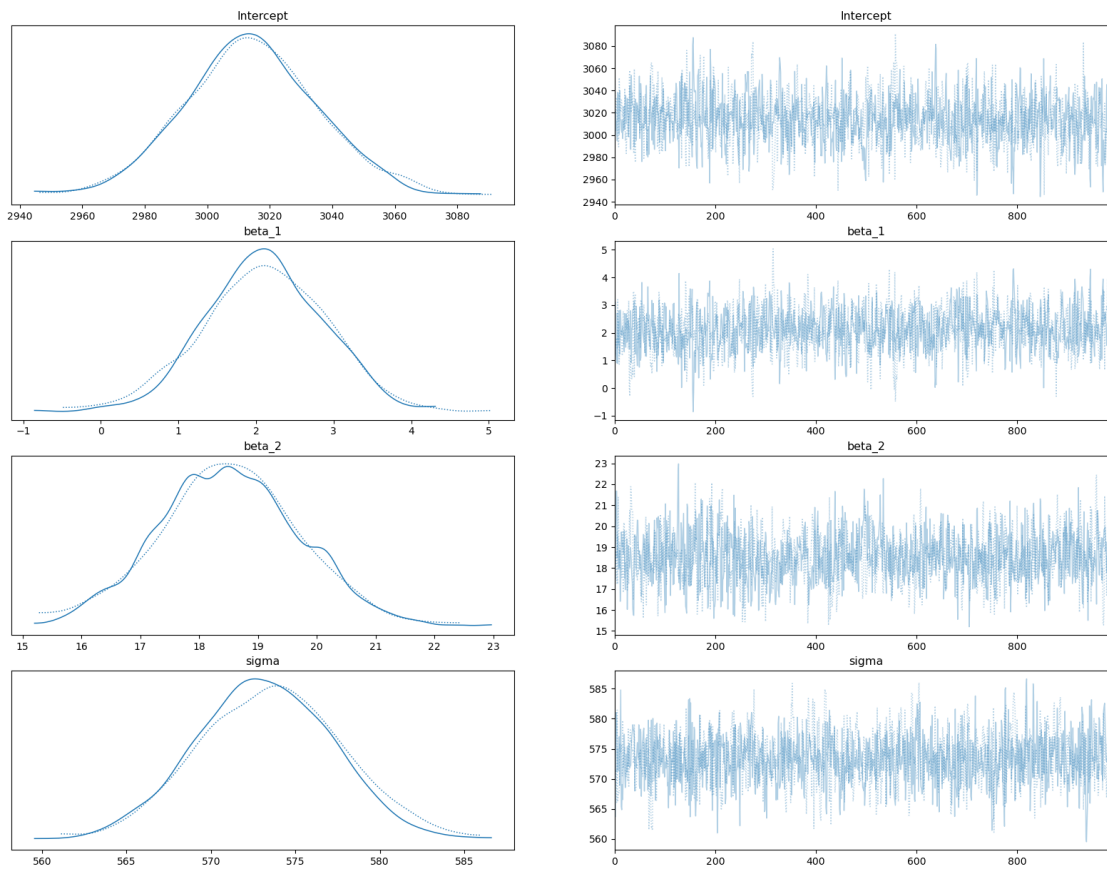
Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 40 seconds.

```
[30]: az.plot_trace(trace, figsize=(20, 15))
```

```

[30]: array([[<AxesSubplot:title={'center': 'Intercept'}>,
<AxesSubplot:title={'center': 'Intercept'}>],
[<AxesSubplot:title={'center': 'beta_1'}>,
<AxesSubplot:title={'center': 'beta_1'}>],
[<AxesSubplot:title={'center': 'beta_2'}>,
<AxesSubplot:title={'center': 'beta_2'}>],
[<AxesSubplot:title={'center': 'sigma'}>,
<AxesSubplot:title={'center': 'sigma'}>]], dtype=object)

```



```
[31]: np.mean(trace.posterior["Intercept"].values)
```

```
[31]: 3014.142456984229
```

```
[32]: np.mean(trace.posterior["beta_1"].values)
```

```
[32]: 2.1058263417649217
```

```
[33]: np.mean(trace.posterior["beta_2"].values)
```

```
[33]: 18.516514695016042
```

We see that the posterior prediction given by the bayesian model is  $W_i = 3015.256 + 2.086M_i + 18.471P_i$ . This is very similar to the one given by the frequentist model. This again means that older mothers have babies with higher birth weights, and mothers who go to more prenatal visits also have babies with higher birth weights.

### 6.4.1 Model Checking

```
[34]: with bayes_model:
      ppc = pm.sample_posterior_predictive(
          trace, var_names=["beta_1", "beta_2", "Intercept", "y"]
      )
```

<IPython.core.display.HTML object>


<IPython.core.display.HTML object>

```
[35]: ppc
```

```
[35]: {'beta_1': array([1.608638 , 3.11911254, 1.25982419, ..., 2.1650516 ,
 2.38930908,
 1.65561139]),
      'beta_2': array([19.25056496, 17.46079249, 19.94211726, ..., 19.70850229,
 17.00464306, 17.67605704]),
      'Intercept': array([3010.62971253, 2999.66959585, 3021.44245843, ...,
 2997.22908972,
 3031.87069906, 3026.07162836]),
      'y': array([[2631.85897825, 2893.73461445, 3145.97074589, ..., 2723.39482088,
 2480.1428695 , 4635.43100947],
 [2981.54851896, 3679.75829088, 3771.45585581, ..., 3418.21215197,
 3699.64485004, 3141.13181795],
 [2845.71825798, 3153.07848564, 2422.53078112, ..., 3773.20936441,
 4007.32236596, 3945.3082004 ],
 ...,
 [4088.58951957, 3954.3746849 , 2068.08594291, ..., 3078.62229034,
 3354.71237771, 3598.06527111],
 [3653.94283063, 3454.13215897, 4321.2614436 , ..., 3852.82097892,
 3483.37595291, 2656.09877589],
 [3401.96189087, 3294.24525591, 3405.46915864, ..., 3055.15488235,
 3489.73823729, 2898.96480928]])}
```


```
[36]: for i in np.arange(100):
      sns.histplot(ppc['y'][i])
      plt.title("PPC of Birth Weights")
```

```
[36]: Text(0.5, 1.0, 'PPC of Birth Weights')
```



```
[37]: #Actual birthweight histogram
sns.histplot(birth['DBWT'])
plt.title("Actual distribution of Birth Weights")
```

```
[37]: Text(0.5, 1.0, 'Actual distribution of Birth Weights')
```




As we can see from the plots above, the PPC samples have a similar distribution as our data, so the model is a reasonable fit for the data.

### 6.4.2 Uncertainty Quantification

```
[38]: az.plot_posterior(trace, ['Intercept', 'beta_1', 'beta_2'], round_to = 3)
```

```
[38]: array([<AxesSubplot:title={'center':'Intercept'}>,
<AxesSubplot:title={'center':'beta_1'}>,
<AxesSubplot:title={'center':'beta_2'}>], dtype=object)
```



We see that the highest density interval (another term for credible interval) for the intercept is [2974, 3053]. The HDI for the coefficient of MAGER is [0.605, 3.41]. The HDI for the coefficient of PREVIS is [16.1, 20.8].